

Figure 4 — Inference Latency per Architecture and Target MCU

